

Institut
québécois
d'intelligence
artificielle



Mila

Humanware project Detection of text in natural scenes

Margaux Luck, PhD
Jeremy Pinto

HumanWare

Over 25 years of service to people with vision loss

Recognized for its capacity to innovate

(Digital reading systems, talking book players, handheld talking GPS)

A worldwide network of distributors

(North America, Europe & Australasia)

150 employees making independence possible

(15% of our employees are also our clients)

#1 worldwide in blindness service

Products supporting over 25 languages

(including the Braille script)

Products available in more than 45 countries



Drummondville, Canada
Head Office
Operations Canada & USA



Rushden, Angleterre
Operations Europe,
Middle East & Africa

Longueuil, Canada
Marketing & R&D



Sydney, Australia
Operations Australasia



Humanware project

- **Task:** Detection of text in natural scenes, with its location and the ability to guide the user to get the full text (alignment) so that the text is interpretable by a recognition engine.
- **Use case:** Detection of house numbers, text on a bus stop sign, or text on a storefront (name, product details, opening hours, ...)
- **Constraints:** Execution time, online vs. offline, memory usage (in the case of a mobile application), etc.

Focus on door number detection



668

- Help blind persons find their way around
- Make sure that's the right house
- (Long term) Facilitate micronavigation

Data

The Street View House Numbers (SVHN) Dataset:

- Open-source
- ~200k street numbers
- bounding boxes and class labels for individual digits, giving about 600k digits total

Limitation:

- zoom on the numbers, lack of background
- no negative examples (i.e. no images without numbers)



[Yuval Netzer, Tao Wang, Adam Coates, Alessandro Bissacco, Bo Wu, Andrew Y. Ng Reading Digits in Natural Images with Unsupervised Feature Learning NIPS Workshop on Deep Learning and Unsupervised Feature Learning 2011.]

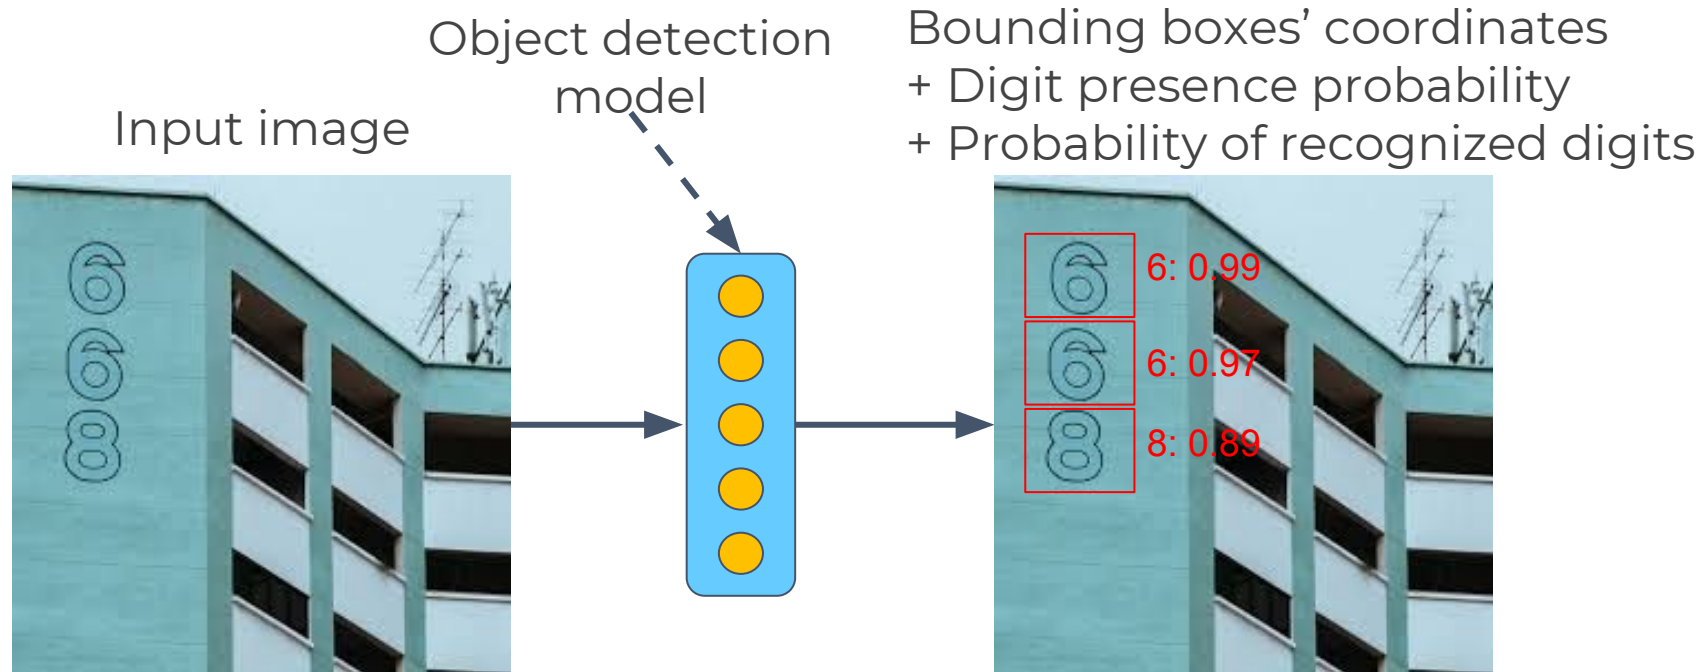
Ongoing data collection

In house dataset

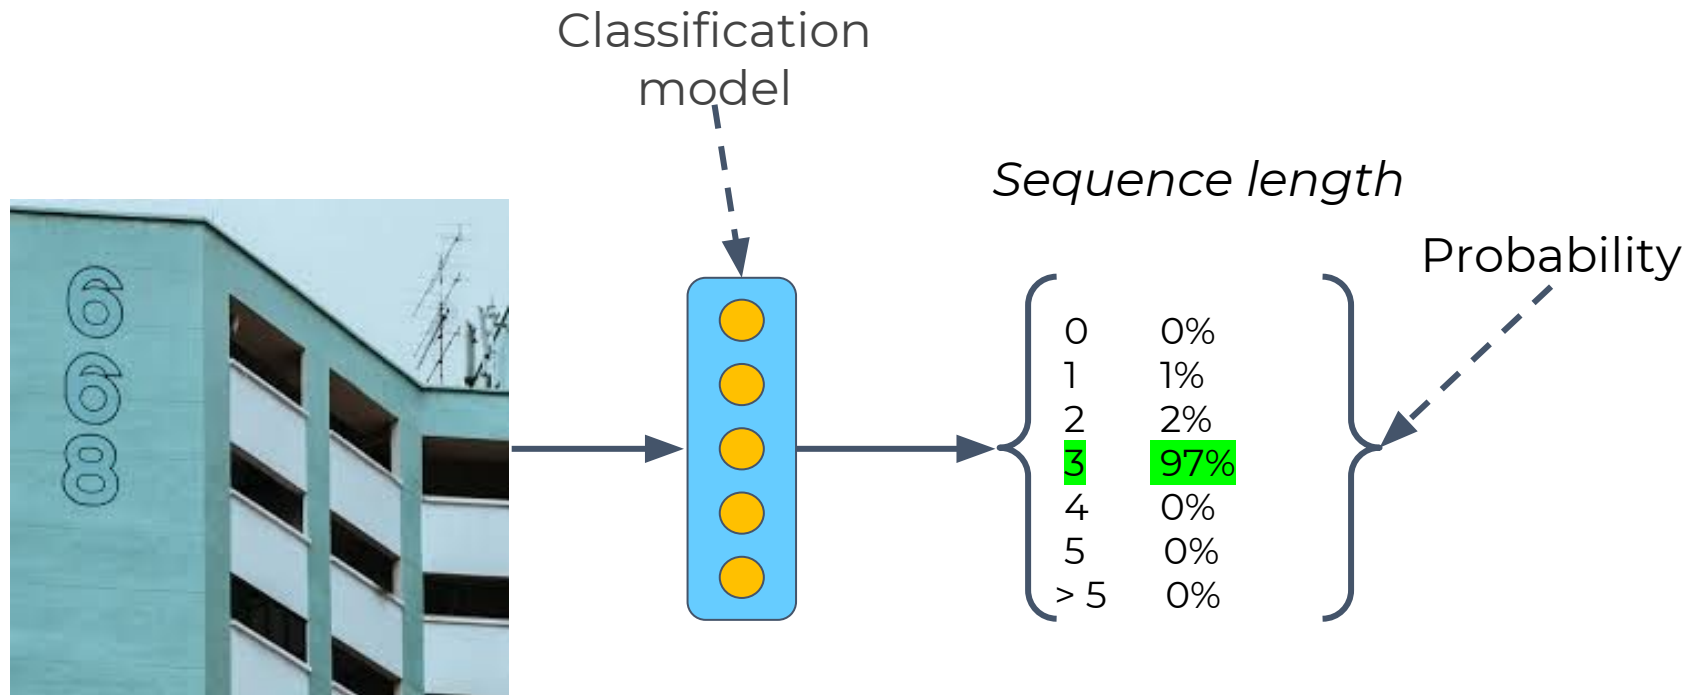
- Aim to be open-source
- Video of street views (case studies defined with the users)
- No ground truth bounding boxes available
- Centroid of the door numbers and sequence label
- The collection started
- You can help!



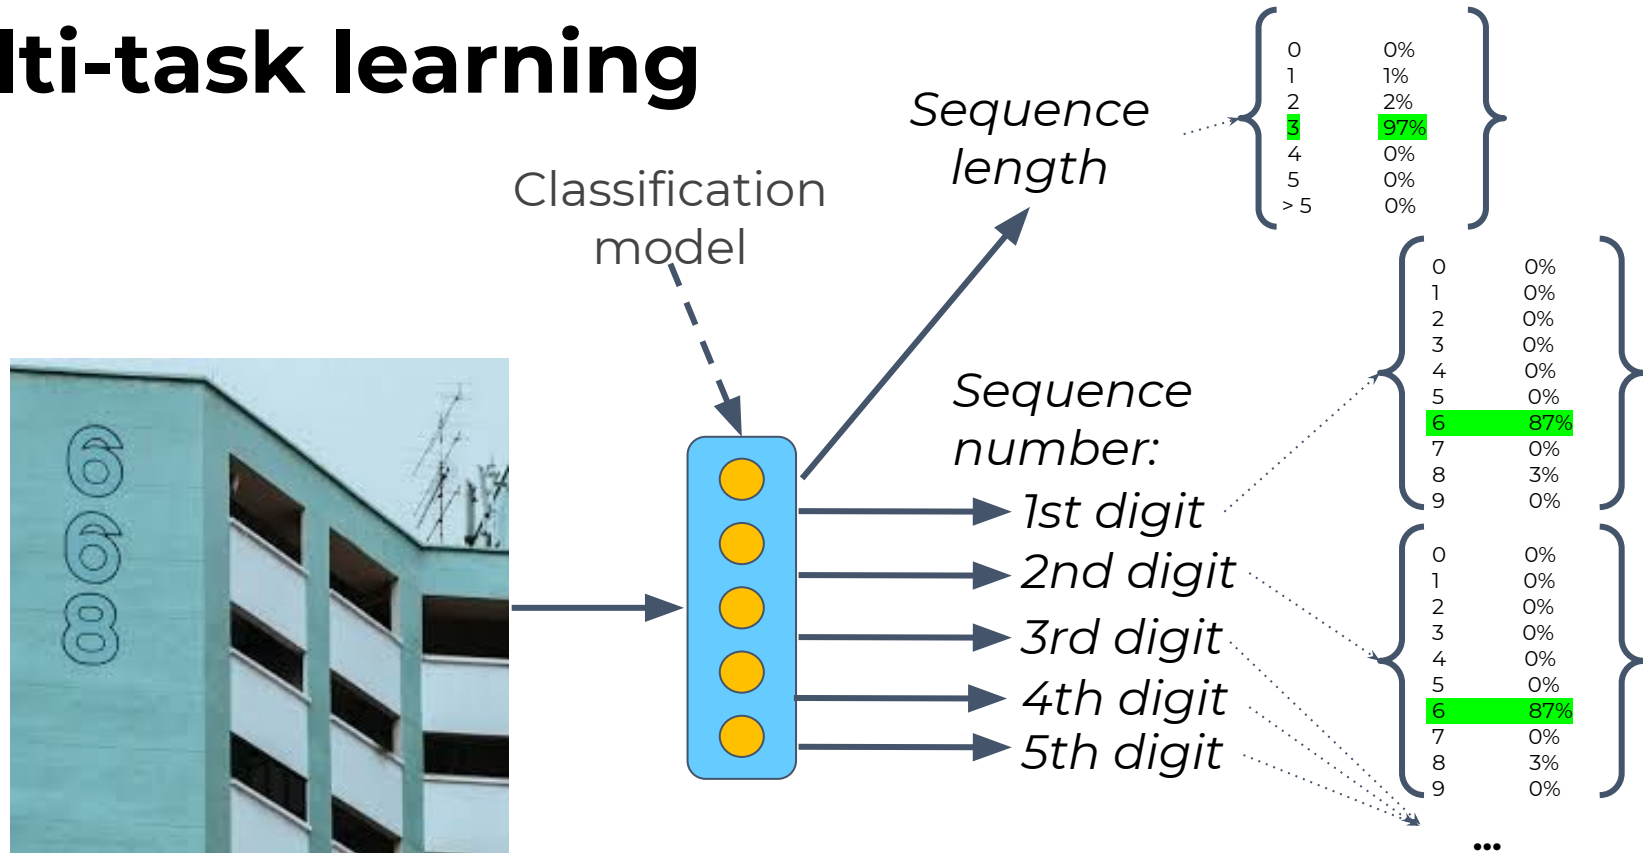
Object detection



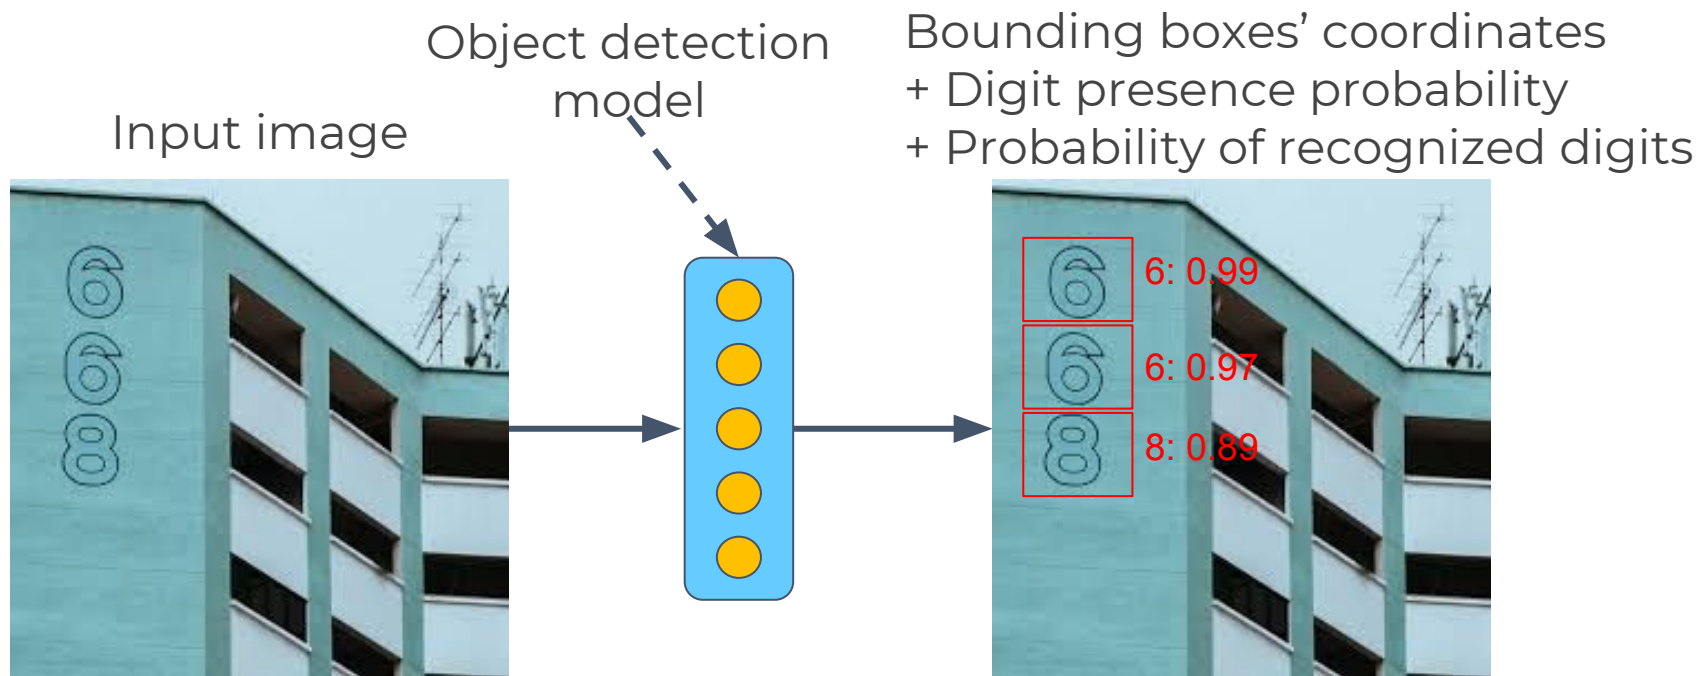
Classification



Multi-task learning



Object detection (multi-task learning)



Official evaluation metrics

- Sequence length accuracy as initial metric

That will be replaced with:

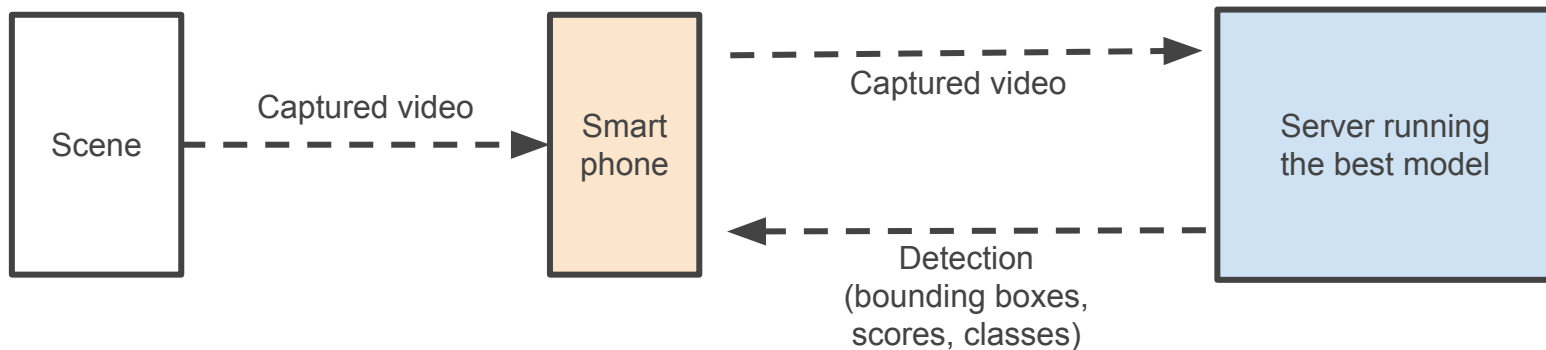
- Sequence transcription accuracy
 - $\frac{\text{\# correct sequences}}{\text{total \# of sequences}}$

Informative evaluation metrics


- Sequence length accuracy
 - Sequence transcription accuracy
 - Digit-level accuracy
 - Coverage at 98% sequence transcription accuracy
 - etc.
- + *For object detection models (that include bounding boxes)*
- Sequence intersection over union (IoU)
 - $\text{IoU} = (\text{area of overlap}) / (\text{area of union})$
 - Digit-level IoU
 - etc.

Demo

Door number detection: server solution



Door number detection: embedded solution



Block 1 instructions / expected timeline

	2019/01/14 week	2019/01/21 week	2019/01/28 week	2019/02/04 week
Tasks / Homework	<ul style="list-style-type: none">• Literature review of available datasets• Read the Goodfellow et al. 2013 article	<ul style="list-style-type: none">• Code the data loader for SVHN as in Goodfellow et al. 2013• Implement data visualization functions	<ul style="list-style-type: none">• Train a CNN for the classification of the sequence length	<ul style="list-style-type: none">• Write a short report summarizing the work, and results• (Peer-) Review of other teams' code
Objectives / Deliverables	<ul style="list-style-type: none">• Justify the choice of using the SVHN dataset and to collect an in-house dataset	<ul style="list-style-type: none">• SVHN data loader	<ul style="list-style-type: none">• CNN for the classification of the sequence length	<ul style="list-style-type: none">• Produce documented code and report summarizing the experimental work• Provide model for blind test set evaluation• Complete the peer code review

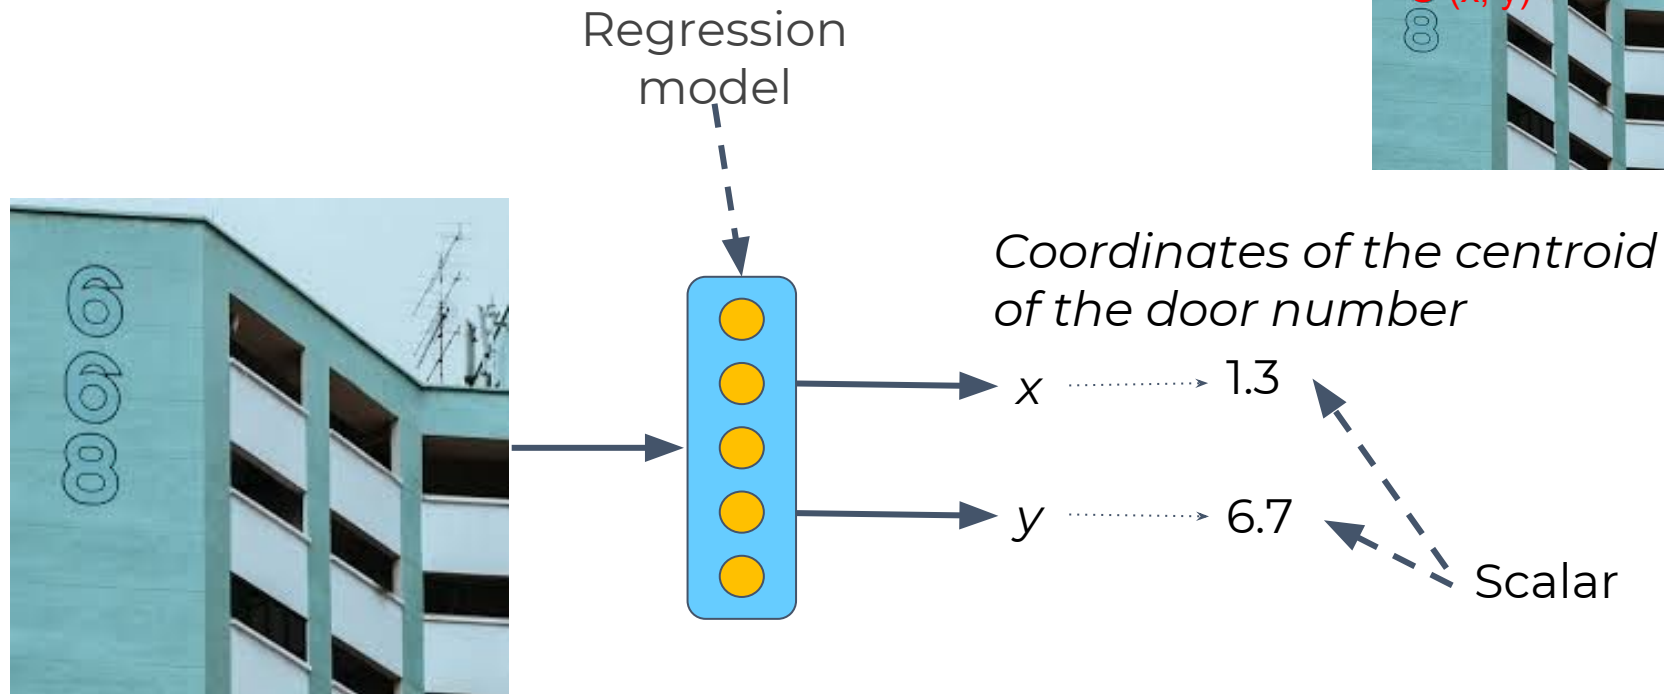
Block 2 instructions / expected timeline

	2019/02/11 week	2019/02/18 week	2019/02/25 week	2019/03/11 week
Tasks / Homework	<ul style="list-style-type: none"> Read and understand the documentation (code) and reports 	<ul style="list-style-type: none"> Implementation of loss to classify the different digits of the image (Multi-Task learning) 	<ul style="list-style-type: none"> implementation + evaluation of different models (AlexNet, VGG, SqueezeNet,...) / inference speed, metrics Search for hyper-parameters, use of TensorboardX, management of various models, comparative performance grids 	<ul style="list-style-type: none"> Search for hyper-parameters, use of TensorboardX, management of various models, comparative performance grids Write a short report summarizing the work, and results (Peer-) Review of other teams' code
Objectives / Deliverables	<ul style="list-style-type: none"> Understand the documentation (code) and reports 	<ul style="list-style-type: none"> Multi-task learning model for the sequence digits recognition 	<ul style="list-style-type: none"> Find the best set of hyper-parameters and visualize the models training, ... 	<ul style="list-style-type: none"> Produce documented code and report summarizing the experimental work Provide model for blind test set evaluation Complete the peer code review

Block 3 instructions / expected timeline

	2019/03/18 week	2019/03/25 week	2019/04/01 week	2019/04/08 week
Tasks / Homework	<ul style="list-style-type: none">• Read and understand the documentation (code) and reports	<ul style="list-style-type: none">• Add the object detection part to the classification model (+ number extraction)• Application and finetuning (i.e. transfer learning) of the model on the Humanware dataset - they are given the dataloader	<ul style="list-style-type: none">• Finetuning of the models and obtaining the results for the final presentation	<ul style="list-style-type: none">• Preparation of the final presentation of 15 min + 5 min of questions which describes the evolution of the project through the 3 blocks• (Peer-) Review of other teams' code
Objectives / Deliverables	<ul style="list-style-type: none">• Understand the documentation (code) and reports	<ul style="list-style-type: none">• Object detection model of the door digit	<ul style="list-style-type: none">• Finetune object detection model	<ul style="list-style-type: none">• Produce documented code and presentation summarizing the experimental work of the 3 blocks• Provide model for blind test set evaluation• Complete the peer code review

Multi-task learning



Multi-task learning

